

6 DETERMINISTIC ILLUSTRATIVE SUCCESS ERRORS - NOT A PREVALENCE ESTIMATE

FP Borderline LONG BUT UNSUCCESSFUL visualwebarena Length and text-volume statistics cannot tell productive evidence gathering from a blocked destination. Semantic evidence needed: yes	FP Median confidence TERMINATION MISMATCH webarena The representation observes activity and terminal-action structure but not whether the user received the requested semantic answer. Semantic evidence needed: yes	FP High confidence LONG BUT UNSUCCESSFUL workarena Repeated terminal-form actions and long answer strings can look completion-like while answer correctness remains unavailable to task-agnostic structure. Semantic evidence needed: yes
FN Borderline SHORT BUT SUCCESSFUL visualwebarena A short efficient path can contain enough semantic evidence even when structural counts resemble under-execution. Semantic evidence needed: yes	FN Median confidence TERMINATION MISMATCH webarena Structural features cannot reconcile task state completion with an incomplete natural-language handoff. Semantic evidence needed: yes	FN High confidence REPETITIVE BUT PROGRESSING webarena Exact structural repetition cannot determine whether repeated submenu operations are stalled or eventually productive. Semantic evidence needed: yes